

## EXPERIMENT NO. 1

**Problem Statement:** Introduction to Data science and Data preparation using Pandas steps.

### Theory:

Data science is a multidisciplinary field focused on deriving valuable insights from both structured and unstructured data through scientific methods, algorithms, and systems. A crucial initial step in any data science project is data preparation, which involves cleaning, transforming, and structuring raw data to improve its quality and ensure it is suitable for analysis.

In this experiment, we implemented data preprocessing techniques using the Pandas library in Python. The dataset analyzed consists of records of car accidents in NYC from 2020, including key attributes such as the number of injuries, fatalities, latitude, longitude, contributing factors, and vehicle types involved. Initially, the dataset contained missing values, inconsistencies, and redundant columns, making it necessary to conduct comprehensive cleaning and preprocessing to enhance data quality.

The following key steps were carried out in this process:

- Importing the dataset into Pandas.
- Detecting and handling missing values.
- Removing redundant columns.
- Applying ordinal encoding to categorical variables.
- Identifying and managing outliers.
- Standardizing and normalizing numerical features.

#### 1. Loading data into pandas:

```
import pandas as pd
df = pd.read_csv(r"C:\Users\bhumil\OneDrive\文档\ds_lab_csv\nyc_accidents.csv")
df.info()
```

```

----- RESTART: C:\Users\bhumi\OneDrive\文档\ds_lab_csv\expl.py -----
--
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 74881 entries, 0 to 74880
Data columns (total 29 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   CRASH DATE                                74881 non-null  object
1   CRASH TIME                                74881 non-null  object
2   BOROUGH                                   49140 non-null  object
3   ZIP CODE                                  49134 non-null  float64
4   LATITUDE                                  68935 non-null  float64
5   LONGITUDE                                 68935 non-null  float64
6   LOCATION                                  68935 non-null  object
7   ON STREET NAME                           55444 non-null  object
8   CROSS STREET NAME                        35681 non-null  object
9   OFF STREET NAME                          19437 non-null  object
10  NUMBER OF PERSONS INJURED                74881 non-null  int64
11  NUMBER OF PERSONS KILLED                 74881 non-null  int64
12  NUMBER OF PEDESTRIANS INJURED            74881 non-null  int64
13  NUMBER OF PEDESTRIANS KILLED             74881 non-null  int64
14  NUMBER OF CYCLIST INJURED                74881 non-null  int64
15  NUMBER OF CYCLIST KILLED                 74881 non-null  int64
16  NUMBER OF MOTORIST INJURED               74881 non-null  int64
17  NUMBER OF MOTORIST KILLED                74881 non-null  int64
18  CONTRIBUTING FACTOR VEHICLE 1             74577 non-null  object
19  CONTRIBUTING FACTOR VEHICLE 2             59285 non-null  object
20  CONTRIBUTING FACTOR VEHICLE 3             6765 non-null   object
21  CONTRIBUTING FACTOR VEHICLE 4             1851 non-null   object
22  CONTRIBUTING FACTOR VEHICLE 5             523 non-null    object
23  COLLISION_ID                             74881 non-null  int64
24  VEHICLE TYPE CODE 1                      74246 non-null  object
25  VEHICLE TYPE CODE 2                      53638 non-null  object
26  VEHICLE TYPE CODE 3                      6424 non-null   object
27  VEHICLE TYPE CODE 4                      1771 non-null   object
28  VEHICLE TYPE CODE 5                      503 non-null    object
dtypes: float64(3), int64(9), object(17)
memory usage: 16.6+ MB

```

## 2. Description of the dataset.

The dataset consists of features and instances regarding Car accidents in NYC in 2020. Major features like the amount of people killed, amount of people injured, latitude, longitude, etc collectively make up this dataset. Starting off, we get to see that there are innumerable null values, missing values, inconsistent data within the dataset. We need to use Pandas in Python to clean and process the data within it.

## 3. Drop columns that are not useful:

```

import pandas as pd
df = pd.read_csv(r"C:\Users\bhumi\OneDrive\文档\ds_lab_csv\nyc_accidents.csv")
cols_to_drop = [
    'CONTRIBUTING FACTOR VEHICLE 3',
    'CONTRIBUTING FACTOR VEHICLE 4',
    'CONTRIBUTING FACTOR VEHICLE 5',
    'VEHICLE TYPE CODE 3',

```

```
'VEHICLE TYPE CODE 4',
'VEHICLE TYPE CODE 5',
'OFF STREET NAME'
]
```

```
df = df.drop(columns=[col for col in cols_to_drop if col in df.columns], axis=1)
```

```
Index: 23959 entries, 0 to 74879
```

```
Data columns (total 22 columns):
```

#	Column	Non-Null Count	Dtype
0	CRASH DATE	23959 non-null	object
1	CRASH TIME	23959 non-null	object
2	BOROUGH	23959 non-null	object
3	ZIP CODE	23954 non-null	float64
4	LATITUDE	23959 non-null	float64
5	LONGITUDE	23959 non-null	float64
6	LOCATION	23959 non-null	object
7	ON STREET NAME	23959 non-null	object
8	CROSS STREET NAME	23950 non-null	object
9	NUMBER OF PERSONS INJURED	23959 non-null	int64
10	NUMBER OF PERSONS KILLED	23959 non-null	int64
11	NUMBER OF PEDESTRIANS INJURED	23959 non-null	int64
12	NUMBER OF PEDESTRIANS KILLED	23959 non-null	int64
13	NUMBER OF CYCLIST INJURED	23959 non-null	int64
14	NUMBER OF CYCLIST KILLED	23959 non-null	int64
15	NUMBER OF MOTORIST INJURED	23959 non-null	int64
16	NUMBER OF MOTORIST KILLED	23959 non-null	int64
17	CONTRIBUTING FACTOR VEHICLE 1	23959 non-null	object
18	CONTRIBUTING FACTOR VEHICLE 2	23734 non-null	object
19	COLLISION_ID	23959 non-null	int64
20	VEHICLE TYPE CODE 1	23959 non-null	object
21	VEHICLE TYPE CODE 2	21723 non-null	object

```
dtypes: float64(3), int64(9), object(10)
```

```
memory usage: 4.2+ MB
```

After saving,

running and viewing our updated dataset, we see that the unnecessary columns have been eliminated.

#### 4. Dropping rows with missing values:

```
df = df.dropna()
```

```
df.info()
```

```
[0 rows x 23 columns]
<class 'pandas.core.frame.DataFrame'>
Index: 0 entries
Data columns (total 23 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   CRASH DATE                                0 non-null      object
1   CRASH TIME                                0 non-null      object
2   BOROUGH                                    0 non-null      object
3   ZIP CODE                                  0 non-null      float64
4   LATITUDE                                  0 non-null      float64
5   LONGITUDE                                 0 non-null      float64
6   LOCATION                                  0 non-null      object
7   ON STREET NAME                           0 non-null      object
8   CROSS STREET NAME                        0 non-null      object
9   OFF STREET NAME                          0 non-null      object
10  NUMBER OF PERSONS INJURED                 0 non-null      int64
11  NUMBER OF PERSONS KILLED                  0 non-null      int64
12  NUMBER OF PEDESTRIANS INJURED             0 non-null      int64
13  NUMBER OF PEDESTRIANS KILLED              0 non-null      int64
14  NUMBER OF CYCLIST INJURED                 0 non-null      int64
15  NUMBER OF CYCLIST KILLED                  0 non-null      int64
16  NUMBER OF MOTORIST INJURED                0 non-null      int64
17  NUMBER OF MOTORIST KILLED                 0 non-null      int64
18  CONTRIBUTING FACTOR VEHICLE 1             0 non-null      object
19  CONTRIBUTING FACTOR VEHICLE 2             0 non-null      object
20  COLLISION_ID                             0 non-null      int64
21  VEHICLE TYPE CODE 1                       0 non-null      object
22  VEHICLE TYPE CODE 2                       0 non-null      object
dtypes: float64(3), int64(9), object(11)
memory usage: 0.0+ bytes
```

The `.dropna()` function, by default, removes any row containing at least one NaN value, which could result in dropping most or all of the rows, especially if several columns have missing data. To address this issue, you can use the `thresh` parameter to specify a minimum number of non-null values required in each row to retain it. By setting `thresh=21`, you ensure that rows with at least 21 non-null values remain in the dataset, while rows with fewer than 21 non-null values are dropped.

```
df = df.dropna(thresh=21)
df.info()
```

```

\Class pandas.core.frame.DataFrame /
Index: 23959 entries, 0 to 74879
Data columns (total 22 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   CRASH DATE                               23959 non-null  object
1   CRASH TIME                               23959 non-null  object
2   BOROUGH                                  23959 non-null  object
3   ZIP CODE                                 23954 non-null  float64
4   LATITUDE                                23959 non-null  float64
5   LONGITUDE                               23959 non-null  float64
6   LOCATION                                 23959 non-null  object
7   ON STREET NAME                           23959 non-null  object
8   CROSS STREET NAME                        23950 non-null  object
9   NUMBER OF PERSONS INJURED                23959 non-null  int64
10  NUMBER OF PERSONS KILLED                 23959 non-null  int64
11  NUMBER OF PEDESTRIANS INJURED            23959 non-null  int64
12  NUMBER OF PEDESTRIANS KILLED             23959 non-null  int64
13  NUMBER OF CYCLIST INJURED                23959 non-null  int64
14  NUMBER OF CYCLIST KILLED                 23959 non-null  int64
15  NUMBER OF MOTORIST INJURED               23959 non-null  int64
16  NUMBER OF MOTORIST KILLED               23959 non-null  int64
17  CONTRIBUTING FACTOR VEHICLE 1            23959 non-null  object
18  CONTRIBUTING FACTOR VEHICLE 2            23734 non-null  object
19  COLLISION_ID                             23959 non-null  int64
20  VEHICLE TYPE CODE 1                     23959 non-null  object
21  VEHICLE TYPE CODE 2                     21723 non-null  object
dtypes: float64(3), int64(9), object(10)

```

## 5. Taking care of missing values:

First we need to find out the number of unique values in each column, so we run

```

unique_counts = df.nunique()
print(unique_counts)

```

Based on the number of unique values, the columns can be categorized in the following:

- Low-Cardinality Categorical Columns - These are columns with very few unique values.
- High-Cardinality Categorical Columns - These are columns with a large number of unique values.

```

CRASH DATE                242
CRASH TIME                1401
BOROUGH                   5
ZIP CODE                  183
LATITUDE                  11723
LONGITUDE                 10820
LOCATION                   12774
ON STREET NAME            2782
CROSS STREET NAME        3335
NUMBER OF PERSONS INJURED  11
NUMBER OF PERSONS KILLED   2
NUMBER OF PEDESTRIANS INJURED  3
NUMBER OF PEDESTRIANS KILLED  2
NUMBER OF CYCLIST INJURED   3
NUMBER OF CYCLIST KILLED    2
NUMBER OF MOTORIST INJURED  11
NUMBER OF MOTORIST KILLED   2
CONTRIBUTING FACTOR VEHICLE 1  52
CONTRIBUTING FACTOR VEHICLE 2  39
COLLISION_ID              23959
VEHICLE TYPE CODE 1        122
VEHICLE TYPE CODE 2        158
dtype: int64

```

Thus, BOROUGH, NUMBER OF PERSONS KILLED, PEDESTRIANS KILLED, CYCLIST KILLED, MOTORIST KILLED, NUMBER OF PEDESTRIANS INJURED, CYCLIST INJURED, MOTORIST INJURED are low-cardinality columns. And ON STREET NAME, CROSS STREET NAME, VEHICLE TYPE CODE 1, VEHICLE TYPE CODE 2, CONTRIBUTING FACTOR VEHICLE 1, CONTRIBUTING FACTOR VEHICLE 2 are high-cardinality columns.

```

low_cardinality_cols = ["BOROUGH"]
df[low_cardinality_cols] =
df[low_cardinality_cols].fillna(df[low_cardinality_cols].mode().iloc[0])

```

```

high_cardinality_cols = ["ON STREET NAME", "CROSS STREET NAME", "VEHICLE
TYPE CODE 1", "VEHICLE TYPE CODE 2", "CONTRIBUTING FACTOR VEHICLE 1",
"CONTRIBUTING FACTOR VEHICLE 2"]
df[high_cardinality_cols] = df[high_cardinality_cols].fillna("Unknown")

```

For, numeric columns like ZIP CODE, LATITUDE & LONGITUDE we do the following

```

df["ZIP CODE"] = df["ZIP CODE"].fillna(df["ZIP CODE"].mode()[0]) df["LATITUDE"] =
df["LATITUDE"].fillna(df["LATITUDE"].median()) df["LONGITUDE"] =
df["LONGITUDE"].fillna(df["LONGITUDE"].median())

```

Thus, number of null values  
`print(df.isnull().sum().sum())`

```
<class 'pandas.core.frame.DataFrame'>
Index: 23959 entries, 0 to 74879
Data columns (total 22 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CRASH DATE                            23959 non-null  object
1   CRASH TIME                            23959 non-null  object
2   BOROUGH                               23959 non-null  object
3   ZIP CODE                              23959 non-null  float64
4   LATITUDE                              23959 non-null  float64
5   LONGITUDE                             23959 non-null  float64
6   LOCATION                              23959 non-null  object
7   ON STREET NAME                        23959 non-null  object
8   CROSS STREET NAME                    23959 non-null  object
9   NUMBER OF PERSONS INJURED            23959 non-null  int64
10  NUMBER OF PERSONS KILLED              23959 non-null  int64
11  NUMBER OF PEDESTRIANS INJURED        23959 non-null  int64
12  NUMBER OF PEDESTRIANS KILLED         23959 non-null  int64
13  NUMBER OF CYCLIST INJURED             23959 non-null  int64
14  NUMBER OF CYCLIST KILLED              23959 non-null  int64
15  NUMBER OF MOTORIST INJURED            23959 non-null  int64
16  NUMBER OF MOTORIST KILLED             23959 non-null  int64
17  CONTRIBUTING FACTOR VEHICLE 1         23959 non-null  object
18  CONTRIBUTING FACTOR VEHICLE 2         23959 non-null  object
19  COLLISION_ID                         23959 non-null  int64
20  VEHICLE TYPE CODE 1                   23959 non-null  object
21  VEHICLE TYPE CODE 2                   23959 non-null  object
dtypes: float64(3), int64(9), object(10)
```

## 6. Creating dummy variables:

# Define the categorical columns you want to encode

```
categorical_columns = [
    'BOROUGH',
    'NUMBER OF PERSONS INJURED',
    'NUMBER OF PERSONS KILLED',
    'NUMBER OF PEDESTRIANS INJURED',
    'NUMBER OF PEDESTRIANS KILLED',
    'NUMBER OF CYCLIST INJURED',
    'NUMBER OF CYCLIST KILLED',
    'NUMBER OF MOTORIST INJURED',
    'NUMBER OF MOTORIST KILLED',
    'CONTRIBUTING FACTOR VEHICLE 1',
```



```
'CONTRIBUTING FACTOR VEHICLE 2'  
]
```

```
# Initialize and apply the encoder
```

```
encoder = OrdinalEncoder(handle_unknown='use_encoded_value', unknown_value=-1)  
df[categorical_columns] = encoder.fit_transform(df[categorical_columns])
```

```
# Ensure there are no missing values before converting to int
```

```
df[categorical_columns] = df[categorical_columns].fillna(-1).astype(int)
```

```
CRASH DATE ... VEHICLE TYPE CODE 2_van  
0 2020-08-29 ... False  
1 2020-08-29 ... False  
8 2020-08-29 ... False  
11 2020-08-29 ... False  
16 2020-08-29 ... False
```

```
[5 rows x 389 columns]
```

## 7. Find out outliers (manually)

In the given dataset, the "NUMBER OF PERSONS INJURED" column contains values that are mostly 0, with a few higher values. The value 15 can be considered an outlier as it is significantly higher compared to the majority of values in this column, which are below 10.

## 8. Applying Standardization

Standardization refers to the technique scaling data to have a mean of 0 and a standard deviation of 1. It ensures that each feature contributes equally to the model without being affected by different scales.

We used **StandardScaler()** from **sklearn.preprocessing** to apply standardization:

**Its effect on our dataset:**

- Transforms numerical values into a standard normal distribution.
- Suitable when data follows a **normal distribution**.
- Useful for models that rely on distance (e.g., KNN, SVM, PCA).

**Mentioned below is the code snippet**

```
# Continuous columns to be standardized or normalized  
continuous_columns = [
```



'LATITUDE', 'LONGITUDE',  
 'NUMBER OF PERSONS INJURED', 'NUMBER OF PERSONS KILLED', 'NUMBER OF PEDESTRIANS INJURED', 'NUMBER OF PEDESTRIANS KILLED', 'NUMBER OF CYCLIST INJURED', 'NUMBER OF CYCLIST KILLED', 'NUMBER OF MOTORIST INJURED', 'NUMBER OF MOTORIST KILLED']

# 1. Standardization (Z-score normalization)

scaler = StandardScaler()

df[continuous\_columns] = scaler.fit\_transform(df[continuous\_columns])

## Applying Normalization:

Normalization scales the data between **0 and 1** by using the minimum and maximum values of each feature.

We applied MinMaxScaler() from sklearn.preprocessing:

## Its effect on our dataset:

- Ensures all values fall within the range [0,1].
- Useful for models that require bounded input (e.g., Neural Networks).
- Prevents large-scale differences between variables from dominating the learning process.

## Dataset before cleaning and processing:

#	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	
1	CRASH DATE	CRASH TIME	BOROUGH	ZIP CODE	LATITUDE	LONGITUDE	LOCATION	ON-STREET	NA CROSS STREET	OFF STREET	NUMBER OF PERSONS INJURED	NUMBER OF PERSONS KILLED	NUMBER OF PEDESTRIANS INJURED	NUMBER OF PEDESTRIANS KILLED	NUMBER OF CYCLIST INJURED	NUMBER OF CYCLIST KILLED	NUMBER OF MOTORIST INJURED	NUMBER OF MOTORIST KILLED
2	2020-08-28	15:40:00	BROOKLYN	10465	40.6821	-73.8336	POINT (-73.8336 40.6821)	AVENUE STRONG AVENUE			0	0	0	0	0	0	0	0
3	2020-08-28	21:00:00	BROOKLYN	11221	40.6806	-73.8184	POINT (-73.8184 40.6806)	AVENUE STRONG AVENUE			2	0	0	0	0	0	0	0
4	2020-08-28	18:20:00			40.8188	-73.8488	POINT (-73.8488 40.8188)	AVENUE STRONG AVENUE			1	0	1	0	0	0	0	0
5	2020-08-28	9:00:00	BROOKLYN	10465	40.6821	-73.8336	POINT (-73.8336 40.6821)			1047 BAYVIEW	0	0	0	0	0	0	0	0
6	2020-08-28	17:10:00	BROOKLYN	11220	40.6489	-73.8338	POINT (-73.8338 40.6489)			4008 BAYVIEW	0	0	0	0	0	0	0	0
7	2020-08-28	2:29:00			40.6821	-73.8488	POINT (-73.8488 40.6821)	AVENUE STRONG AVENUE			1	0	0	0	0	0	0	0
8	2020-08-28	19:30:00	BROOKLYN	10465	40.6821	-73.8336	POINT (-73.8336 40.6821)	AVENUE STRONG AVENUE			0	0	0	0	0	0	0	0
9	2020-08-28	9:00:00			40.6806	-73.8184	POINT (-73.8184 40.6806)	AVENUE STRONG AVENUE			0	0	0	0	0	0	0	0
10	2020-08-28	19:50:00	BROOKLYN	10465	40.6821	-73.8336	POINT (-73.8336 40.6821)	AVENUE STRONG AVENUE			0	0	0	0	0	0	0	0
11	2020-08-28	9:20:00	QUEENSBORO	11365	40.7035	-73.8088	POINT (-73.8088 40.7035)	AVENUE STRONG AVENUE			0	0	0	0	0	0	0	0
12	2020-08-28	9:07:00	QUEENSBORO	11365	40.6821	-73.7874	POINT (-73.7874 40.6821)	AVENUE STRONG AVENUE			0	0	0	0	0	0	0	0
13	2020-08-28	14:00:00	QUEENSBORO	11365	40.7042	-73.7884	POINT (-73.7884 40.7042)	AVENUE STRONG AVENUE			0	0	0	0	0	0	0	0
14	2020-08-28	21:33:00	BROOKLYN	10465	40.6821	-73.8336	POINT (-73.8336 40.6821)	AVENUE STRONG AVENUE			1	0	1	0	0	0	0	0
15	2020-08-28	22:30:00	BROOKLYN	11249	40.7036	-73.8184	POINT (-73.8184 40.7036)	AVENUE STRONG AVENUE			0	0	0	0	0	0	0	0
16	2020-08-28	4:14:00			40.6821	-73.8488	POINT (-73.8488 40.6821)	AVENUE STRONG AVENUE			1	0	0	0	0	0	0	0
17	2020-08-28	9:35:00			40.6806	-73.7384	POINT (-73.7384 40.6806)	AVENUE STRONG AVENUE			0	0	0	0	0	0	0	0
18	2020-08-28	19:50:00	BROOKLYN	11220	40.6489	-73.8338	POINT (-73.8338 40.6489)	AVENUE STRONG AVENUE			0	0	0	0	0	0	0	0
19	2020-08-28	19:50:00	QUEENSBORO	11365	40.7122	-73.8088	POINT (-73.8088 40.7122)	AVENUE STRONG AVENUE			2	0	0	0	0	0	0	0
20	2020-08-28	12:29:00	BROOKLYN	10465	40.6821	-73.8336	POINT (-73.8336 40.6821)	AVENUE STRONG AVENUE			2	0	0	0	0	0	0	0
21	2020-08-28	19:30:00	BROOKLYN	11210	40.7188	-73.8184	POINT (-73.8184 40.7188)	AVENUE STRONG AVENUE			1	0	0	0	0	0	0	0
22	2020-08-28	19:30:00	BROOKLYN	11220	40.6489	-73.8338	POINT (-73.8338 40.6489)	AVENUE STRONG AVENUE			1	0	0	0	0	0	0	0
23	2020-08-28	9:30:00			40.6806	-73.7884	POINT (-73.7884 40.6806)	AVENUE STRONG AVENUE			0	0	0	0	0	0	0	0
24	2020-08-28	9:30:00			40.6821	-73.7336	POINT (-73.7336 40.6821)	AVENUE STRONG AVENUE			0	0	0	0	0	0	0	0
25	2020-08-28	19:30:00			40.6806	-73.8336	POINT (-73.8336 40.6806)	AVENUE STRONG AVENUE			1	0	0	0	0	0	0	0
26	2020-08-28	1:45:00	MANHATTAN	10029	40.7547	-73.8341	POINT (-73.8341 40.7547)	AVENUE STRONG AVENUE			1	0	0	0	0	0	0	0
27	2020-08-28	9:45:00	QUEENSBORO	11365	40.7042	-73.7488	POINT (-73.7488 40.7042)	AVENUE STRONG AVENUE			0	0	0	0	0	0	0	0
28	2020-08-28	9:45:00	QUEENSBORO	11365	40.7042	-73.8184	POINT (-73.8184 40.7042)	AVENUE STRONG AVENUE			0	0	0	0	0	0	0	0

**Conclusion:** The experiment involved cleaning and preprocessing a dataset of NYC car accidents from 2020 using Pandas. Initially, the dataset contained missing values, redundant columns, and categorical data requiring transformation for effective analysis. To address these issues, data cleaning techniques were applied, including the removal of columns with a high percentage of missing values and iterating out incomplete rows using a threshold-based approach. Categorical variables were transformed through ordinal encoding to convert text into numerical values, ensuring consistency. Numerical features were then standardized with StandardScaler to achieve a mean of 0 and a standard deviation of 1, followed by normalization with MinMaxScaler to scale values between 0 and 1. These transformations refined the dataset, eliminating inconsistencies and preparing it for accurate and reliable analysis.